

Laboratory 3: Report

Piotr Harmuszkiewicz

Juan Salmerón Moya

In this laboratory, we try different methods for avoiding memory page faults so programs can run faster in a CUDA environment.

The first method is memory prefetching. By prefetching data (using the `cudaMemPrefetchAsync` function) we can drastically reduce the time our programs take as we evade a lot of page faults. By prefetching, we store in memory the next consecutive pages of the ones we initialize with. The program `vector_add_standard` loses 149ms in memory page faults. In `vector_add_prefetch_gpu` this time is not lost, so vector adding takes 2ms instead of 151ms. The method `cudaDeviceSynchronize` takes less time, and Host to Device takes 34ms instead of 44ms.

The second method is to use strides when initializing data. Instead of loading data in a sequential fashion, we can do it by a grid at a time. By doing this, we eliminate overhead from scheduling and retiring blocks. There can be considerable efficiency gains in simple kernels by doing so. The `initWith` takes 2ms of GPU time in `vector_add_prefetch_gpu_init_gpu` but by doing so, `cudaDeviceSynchronize` and `cudaMemPrefetchAsync` take even less time, falling from around 22ms each to 5 or 3 ms.